

# CatBoost vs. Light GBM vs. XGBoost

수요일, 2월 20, 2019 5:21 오후

클리핑 출처: <https://towardsdatascience.com/catboost-vs-light-gbm-vs-xgboost-5f93620723db>



Towards Data Science

Sign in

Get started

HOME DATA SCIENCE MACHINE LEARNING PROGRAMMING VISUALIZATION AI JOURNALISM PICKS | CONTRIBUTE Q



## CatBoost vs. Light GBM vs. XGBoost



Alvira Swalin [Follow](#)

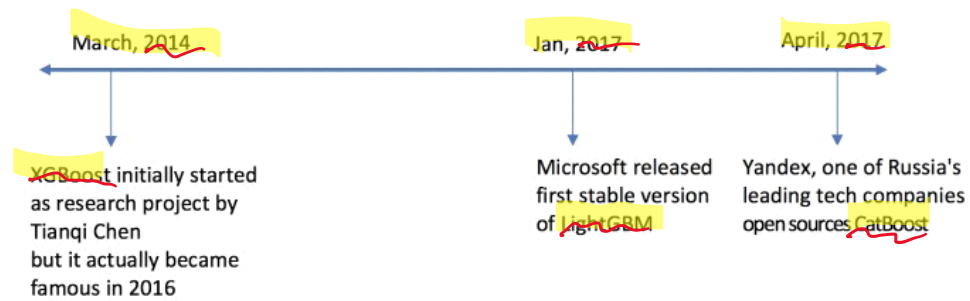
Mar 13, 2018 · 8 min read

Who is going to win this war of predictions and on what cost? Let's explore.



I recently participated in this Kaggle competition (WIDS Datathon by Stanford) where I was able to land up in Top 10 using various boosting algorithms. Since then, I have been very curious about the fine workings of each model including parameter tuning, pros and cons and hence decided to write this blog. Despite the recent re-emergence and popularity of neural networks, I am focusing on boosting algorithms because they are still more useful in the regime of limited training data, little training time and little expertise for parameter tuning.

expertise for parameter tuning.



Since XGBoost (often called GBM Killer) has been in the machine learning world for a longer time now with lots of articles dedicated to it, this post will focus more on CatBoost & LGBM. Below are the topics we will cover-

- Structural Differences
- Treatment of categorical variables by each algorithm
- Understanding Parameters
- Implementation on Dataset
- Performance of each algorithm

## Structural Differences in LightGBM & XGBoost

LightGBM uses a novel technique of Gradient-based One-Side Sampling (GOSS) to filter out the data instances for finding a split value while XGBoost uses pre-sorted algorithm & Histogram-based algorithm for computing the best split. Here instances means observations/samples.

First let us understand how pre-sorting splitting works-

*LGBM = GOSS*

- For each node, enumerate over all features
- For each feature, sort the instances by feature value
- Use a linear scan to decide the best split along that feature basis information gain
- Take the best split solution along all the features

In simple terms, Histogram-based algorithm splits all the data points for a feature into discrete bins and uses these bins to find the split value of histogram. While, it is efficient than pre-sorted algorithm in training speed which enumerates all possible split points on the pre-sorted feature values, it is still behind GOSS in terms of speed.

it is still behind GOSS in terms of speed.

### So what makes this GOSS method efficient?


In AdaBoost, the sample weight serves as a good indicator for the importance of samples. However, in Gradient Boosting Decision Tree (GBDT), there are no native sample weights, and thus the sampling methods proposed for AdaBoost cannot be directly applied. Here comes gradient-based sampling.

Gradient represents the slope of the tangent of the loss function, so logically if gradient of data points are large in some sense, these points are important for finding the optimal split point as they have higher error

GOSS keeps all the instances with large gradients and performs random sampling on the instances with small gradients. For example, let's say I have 500K rows of data where 10k rows have higher gradients. So my algorithm will choose (10k rows of higher gradient+ x% of remaining 490k rows chosen randomly). Assuming x is 10%, total rows selected are 59k out of 500K on the basis of which split value is found.

The basic assumption taken here is that samples with training instances with small gradients have smaller training error and it is already well-trained.

In order to keep the same data distribution, when computing the information gain, GOSS introduces a constant multiplier for the data instances with small gradients. Thus, GOSS achieves a good balance between reducing the number of data instances and keeping the accuracy for learned decision trees.



Leaf with higher gradient/error is used for growing further in LGBM

## How each model treats Categorical Variables?

### CatBoost

CatBoost has the flexibility of giving indices of categorical columns so that



## CatBoost

CatBoost has the flexibility of giving indices of categorical columns so that it can be encoded as one-hot encoding using `one_hot_max_size` (Use one-hot encoding for all features with number of different values less than or equal to the given parameter value).

If you don't pass anything in `cat_features` argument, CatBoost will treat all the columns as numerical variables.

Note: If a column having string values is not provided in the `cat_features`, CatBoost throws an error. Also, a column having default int type will be treated as numeric by default, one has to specify it in `cat_features` to make the algorithm treat it as categorical.

```
from catboost import CatBoostRegressor
# Initialize data
cat_features = [0,1,2]
train_data = [{"a", "b", 1, 4, 5, 6}, {"a", "b", 4, 5, 6, 7}, {"c", "d", 30, 40, 50, 60}]
test_data = [{"a", "b", 2, 4, 6, 8}, {"a", "d", 1, 4, 50, 60}]
train_labels = [10, 20, 30]
# Initialize CatBoostRegressor
model = CatBoostRegressor(iterations=2, learning_rate=1, depth=2)
# Fit model
model.fit(train_data, train_labels, cat_features)
```

For remaining categorical columns which have unique number of categories greater than `one_hot_max_size`, CatBoost uses an efficient method of encoding which is similar to mean encoding but reduces overfitting. The process goes like this—

1. Permuting the set of input observations in a random order. Multiple random permutations are generated
2. Converting the label value from a floating point or category to an integer
3. All categorical feature values are transformed to numeric values using the following formula:

$$avg\_target = \frac{countInClass + prior}{totalCount + 1}$$

Where, `CountInClass` is how many times the label value was equal to "1" for objects with the current categorical feature

value

Prior is the preliminary value for the numerator. It is determined by the starting parameters. `TotalCount` is the total number of objects (up to the current one) that have a categorical feature value matching the current one.

Mathematically, this can be represented using below equation:

Mathematically, this can be represented using below equation:

Let  $\sigma = (\sigma_1, \dots, \sigma_n)$  be the permutation, then  $x_{\sigma_p, k}$  is substituted with

$$\frac{\sum_{j=1}^{p-1} [x_{\sigma_j, k} = x_{\sigma_p, k}] Y_{\sigma_j} + a \cdot P}{\sum_{j=1}^{p-1} [x_{\sigma_j, k} = x_{\sigma_p, k}] + a}, \quad (1)$$

## LightGBM

Similar to CatBoost, LightGBM can also handle categorical features by taking the input of feature names. It does not convert to one-hot coding, and is much faster than one-hot coding. LGBM uses a special algorithm to find the split value of categorical features [\[Link\]](#).

Specific feature names and categorical features:

```
train_data = lgb.Dataset(data, label=label, feature_name=['c1', 'c2', 'c3'],
    categorical_feature=['c3'])
```

Note: You should convert your categorical features to int type before you construct Dataset for LGBM. It does not accept string values even if you passes it through categorical\_feature parameter.

## XGBoost

Unlike CatBoost or LGBM, XGBoost cannot handle categorical features by itself, it only accepts numerical values similar to Random Forest. Therefore one has to perform various encodings like label encoding, mean encoding or one-hot encoding before supplying categorical data to XGBoost.

## Similarity in Hyperparameters

All these models have lots of parameters to tune but we will cover only the important ones. Below is the list of these parameters according to their function and their counterparts across different models.

Function	XGBoost	CatBoost	Light GBM
Important parameters which control overfitting	<ol style="list-style-type: none"> <li><b>learning_rate</b> or <b>eta</b> – optimal values lie between 0.01-0.2</li> <li><b>max_depth</b></li> <li><b>min_child_weight</b>: similar to min_child leaf; default is 1</li> </ol>	<ol style="list-style-type: none"> <li><b>Learning_rate</b></li> <li><b>Depth</b> - value can be any integer up to 16. Recommended - [1 to 10]</li> <li>No such feature like min_child_weight</li> <li><b>l2-leaf-reg</b>: L2 regularization coefficient. Used for leaf value calculation (any positive integer allowed)</li> </ol>	<ol style="list-style-type: none"> <li><b>learning_rate</b></li> <li><b>max_depth</b>: default is 20. Important to note that tree still grows leaf-wise. Hence it is important to tune <b>num_leaves</b> (number of leaves in a tree) which should be smaller than <math>2^{(\text{max\_depth})}</math>. It is a very important parameter for LGBM</li> <li><b>min_data_in_leaf</b>: default=20, alias= min_data, min_child_samples</li> </ol>
Parameters	Not Available	1. <b>cat features</b> : It denotes the	1. <b>categorical feature</b> : specify the

		integer allowed)	alias= min_data, min_child_samples
Parameters for categorical values	Not Available	<ol style="list-style-type: none"> <li>1. <b>cat_features</b>: It denotes the index of categorical features</li> <li>2. <b>one_hot_max_size</b>: Use one-hot encoding for all features with number of different values less than or equal to the given parameter value (max – 255)</li> </ol>	<ol style="list-style-type: none"> <li>1. <b>categorical_feature</b>: specify the categorical features we want to use for training our model</li> </ol>
Parameters for controlling speed	<ol style="list-style-type: none"> <li>1. <b>colsample_bytree</b>: subsample ratio of columns</li> <li>2. <b>subsample</b>: subsample ratio of the training instance</li> <li>3. <b>n_estimators</b>: maximum number of decision trees; high value can lead to overfitting</li> </ol>	<ol style="list-style-type: none"> <li>1. <b>rsm</b>: Random subspace method. The percentage of features to use at each split selection</li> <li>2. No such parameter to subset data</li> <li>3. <b>iterations</b>: maximum number of trees that can be built; high value can lead to overfitting</li> </ol>	<ol style="list-style-type: none"> <li>1. <b>feature_fraction</b>: fraction of features to be taken for each iteration</li> <li>2. <b>bagging_fraction</b>: data to be used for each iteration and is generally used to speed up the training and avoid overfitting</li> <li>3. <b>num_iterations</b>: number of boosting iterations to be performed; default=100</li> </ol>

## Implementation on a Dataset

I am using the Kaggle [Dataset](#) of flight delays for the year 2015 as it has both categorical and numerical features. With approximately 5 million rows, this dataset will be good for judging the performance in terms of both speed and accuracy of tuned models for each type of boosting. I will be using a 10% subset of this data ~ 500k rows.

Below are the features used for modeling:

- MONTH, DAY, DAY\_OF\_WEEK: data type int
- AIRLINE and FLIGHT\_NUMBER: data type int
- ORIGIN\_AIRPORT and DESTINATION\_AIRPORT: data type string
- DEPARTURE\_TIME: data type float
- ARRIVAL\_DELAY: this will be the target and is transformed into boolean variable indicating delay of more than 10 minutes
- DISTANCE and AIR\_TIME: data type float

## **XGBoost**

## **Light GBM**

## CatBoost

While tuning parameters for CatBoost, it is difficult to pass indices for categorical features. Therefore, I have tuned parameters without passing categorical features and evaluated two model—one with and other without categorical features. I have separately tuned `one_hot_max_size` because it does not impact the other parameters.

## Results

XGBoost

Light BGM

CatBoost



	XGBoost	Light BGM		CatBoost	
Parameters Used	max_depth: 50 learning_rate: 0.16 min_child_weight: 1 n_estimators: 200	max_depth: 50 learning_rate: 0.1 num_leaves: 900 n_estimators: 300		depth: 10 learning_rate: 0.15 l2_leaf_reg= 9 iterations: 500 one_hot_max_size = 50	
Training AUC Score	0.999	Without passing indices of categorical features	Passing indices of categorical features	Without passing indices of categorical features	Passing indices of categorical features
		0.992	0.999	0.842	0.887
Test AUC Score	0.789	0.785	0.772	0.752	0.816
Training Time	970 secs	153 secs	326 secs	180 secs	390 secs
Prediction Time	184 secs	40 secs	156 secs	2 secs	14 secs
Parameter Tuning Time (for 81 fits, 200 iteration)	500 minutes	200 minutes		120 minutes	

Top highlight

## End Notes

For evaluating model, we should look into the performance of model in terms of both speed and accuracy.

Keeping that in mind, CatBoost comes out as the winner with maximum accuracy on test set (0.816), minimum overfitting (both train and test accuracy are close) and minimum prediction time & tuning time. But this happened only because we considered categorical variables and tuned one\_hot\_max\_size. If we don't take advantage of these features of CatBoost, it turned out to be the worst performer with just 0.752 accuracy. Hence we learnt that CatBoost performs well only when we have categorical variables in the data and we properly tune them.

Our next performer was XGBoost which generally works well. It's accuracy was quite close to CatBoost even after ignoring the fact that we have categorical variables in the data which we had converted into numerical values for its consumption. However, the only problem with XGBoost is that it is too slow. It was really frustrating to tune its parameters especially (took me 6 hours to run GridSearchCV—very bad idea!). The better way is to tune parameters separately rather than using GridSearchCV. Check out this [blog](#) post to understand how to tune parameters smartly.

Finally, the last place goes to Light GBM. An important thing to note here is that it performed poorly in terms of both speed and accuracy when cat\_features is used. I believe the reason why it performed badly was

that it performed poorly in terms of both speed and accuracy when `cat_features` is used. I believe the reason why it performed badly was because it uses some kind of modified mean encoding for categorical data which caused overfitting (train accuracy is quite high—0.999 compared to test accuracy). However if we use it normally like XGBoost, it can achieve similar (if not higher) accuracy with much faster speed compared to XGBoost

(LGBM—0.785, XGBoost—0.789).

Lastly, I have to say that these observations are true for this particular dataset and may or may not remain valid for other datasets. However, one thing which is true in general is that XGBoost is slower than the other two algorithms.

So which one is your favorite? Please comment with the reasons.  
Any feedback or suggestions for improvement will be really appreciated!

Check out my other blogs [here](#)!

LinkedIn: [www.linkedin.com/in/alvira-swalin](http://www.linkedin.com/in/alvira-swalin)

## Resources

1. [http://learningsys.org/nips17/assets/papers/paper\\_11.pdf](http://learningsys.org/nips17/assets/papers/paper_11.pdf)
2. <https://papers.nips.cc/paper/6907-lightgbm-a-highly-efficient-gradient-boosting-decision-tree.pdf>
3. <https://arxiv.org/pdf/1603.02754.pdf>
4. <https://github.com/Microsoft/LightGBM>
5. <https://www.analyticsvidhya.com/blog/2017/06/which-algorithm-takes-the-crown-light-gbm-vs-xgboost/>
6. <https://stats.stackexchange.com/questions/307555/mathematical-differences-between-gbm-xgboost-lightgbm-catboost>

Machine Learning

Boosting

Data Science

Xgboost

Towards Data Science



3.6K claps



23



**Alvira Swalin**

Data Scientist @Uber,  
MSDS @USF, IIT  
Bombay Alumnus,  
[www.linkedin.com/in/alvira-swalin](http://www.linkedin.com/in/alvira-swalin)

Follow



**Towards Data Science**

Sharing concepts,  
ideas, and codes.

Follow



More from Towards Data Science

## How It Feels to Learn Data Science in 2019



Thomas Nield  
Feb 4 · 16 min r



5.1K



More from Towards Data Science


## What no one will tell you about data science job...



Edouard Harris  
Feb 11 · 6 min



1.94K



More from Towards Data Science

## 10 Python File System Methods You Should Know



Jeff Hale  
Feb 15 · 7 min re



484



### Responses

Write a response...

Conversation between Hugo Lopes and Alvira Swalin.



Hugo Lopes  
Mar 21, 2018 · 1 min read

Hi Alvira,

Great post on a very interesting subject :) Two points:

1. From what I see you didn't do any cross-validation (or am I wrong?), which means that the results might change if you initialize your `train_test_split` with a different `random_seed`. When we are talking about small differences in the

Read more...



30

1 response



Alvira Swalin  
Mar 21, 2018

Hi James,

Thank you for your suggestions. I have used cross-validation to tune hyper-parameters but haven't reported the CV error, so your point makes sense.

2. <https://scikit-optimize.github.io/> looks good. I will check it out for future work.



9



Conversation with Alvira Swalin.

Conversation with Alvira Swalin.



David Olmo Pérez  
Mar 29, 2018

This review was superb! Congratulations for the good work

For me CatBoost has been performing very well recently and I've fully switched to it!

I'm a bit surprised that Hyperparameter tuning turned out to be faster in CatBoost on this review. In my experience LightGBM is the best one (by far) in terms of speed



13

1 response



Alvira Swalin  
Mar 29, 2018 · 1 min read

Thanks David!

Yeah, I agree lightgbm is faster compared to any other method (clearly visible from the training time). I think the reason why parameter tuning took more time is because there were different number of iterations used in each method to get a decent level of accuracy. There is a possibility that I have used more!

[Read more...](#)



2



Conversation with Alvira Swalin.



Chenchao Zang  
Mar 14, 2018

Great Blog. Just to point a name error. The original author of XgBoost is Tianqi Chen.



5

1 response



Alvira Swalin  
Mar 15, 2018

Thanks! I have updated it.



6



Conversation with Alvira Swalin.



Trent Whiting  
Mar 15, 2018

Nice post, hadn't heard of Catboost, definitely check it out, speed and accepting categorical features a nice plus.



Nice post, hadn't heard of Catboost, definitely check it out, speed and accepting categorical features a nice plus.

You should also check out randomized grid search in sklearn for hyperparameter tuning as its generally much more efficient and quicker on new problems.

Thanks for post.



4

1 response 



Alvira Swalin

Mar 15, 2018

Thanks Trent!

Using Random Search is a good idea but the only problem is that it does not give global optima.



3



---

Conversation with Alvira Swalin.



Kesava B

Mar 15, 2018

Hello Alvira,

Nice post. The link to the completion

<https://www.kaggle.com/c/wids2018datathon> is not working. Please check if this need to be updated.



1 response 



Alvira Swalin

Mar 18, 2018

Hi Kesava, thanks for pointing out. I just realized that this competition is visible to only participants. Anyways, I have updated it.



---

Conversation with Alvira Swalin.



Shubhrangshu

May 12, 2018 · 1 min read

Hi,

Very lucid explanation. Great post. I have 2 points relevant to this post:

1. You have clearly put the disclaimer that the result-set given is on the basis of the particular Kaggle-challenge data-set. In this context the benchmark (<https://catboost.yandex/#benchmark>) as put by the Yandex(the

benchmark (<https://catboost.yandex/#benchmark>) as put by the Yandex(the

Read more...



21

2 responses



Alvira Swalin  
May 14, 2018

Hi Shubhrangshu,

Thank you for your feedback. The benchmark by Yandex looks great.

I never heard of Hyperband technique but now I will try to use it.

Thanks.



1



Show all responses



Never miss a story from **Towards Data Science**, when you sign up for Medium. [Learn more](#)

GET UPDATES